

Sorting

1 Insertion Sort

Insertion Sort solves the sorting problem:

Input: A sequence of n numbers $\langle a_1, a_2, \dots, a_n \rangle$

Output: A permutation of the input sequence $\langle a'_1, a'_2, \dots, a'_n \rangle$ where $a'_1 \leq a'_2 \leq \dots \leq a'_n$.

The numbers to be sorted are known as *keys*. These keys are often associated with other data, which we call *satellite data*. A key and satellite data form a *record*. When sorting with respect to some key, the satellite data is also sorted so that it stays with the associated key.

In this text we shall write algorithms in pseudocode. Real implementations in various languages can be found in the same directory as this text. Some notes regarding conventions. We shall organise compound data into *objects*, composed of *attributes*. Parameters will be passed *by value*; that is to say, the called procedure receives its own copy of the parameter. Assignments to this parameter within the called procedure will not be visible to the caller. However, if the parameter is a pointer (for example, pointer x with some attribute f), assignment $x.f = 3$ will be visible to the caller, as will changes to an array, which is passed by a pointer to the first element.

Assume boolean operators are *short circuiting*. That is, for operators *and* and *or*, the left expression will be evaluated first. The right expression will only be evaluated if the full expression's value isn't determined by the left expression. For example, if x *and* y , y won't be evaluated if x is false. Similarly, if x *or* y , y won't be evaluated if x is true.

Insertion sort is an efficient algorithm when sorting a small number of elements.

```

INSERTION-SORT(A, n)
  for i = 2 to n
    key = A[i]
    // Insert A[i] into the sorted subarray A[1:i-1]
    j = i - 1
    while j > 0 and A[j] > key
      A[j + 1] = A[j]
      j = j - 1
    A[j + 1] = key

```

In words, we separate the array of elements into two *subarrays* (that is, contiguous portions of the array), where the left subarray is sorted and the right may not be. We then iterate through the right array, taking the leftmost element in the right array (call this the key) and trying to find its correct position by comparing its value with each element in the left array, from right to left. When there is a value that is less than the key, insert the key to the right of this value, shifting all the values above this along by one. Repeating this with every element in the right array will eventually result in a sorted array. See Figure 1 for a visual representation.




Figure 1: Visual representation of the insertion sort algorithm.

We want to prove that this is a correct algorithm. We can do this by mathematical induction using a *loop invariant*. A loop invariant is a property of a program that is true before and after each iteration of a loop. Thus, if we can formulate the desired outcome of our algorithm (in this case, for the array to be sorted) as a loop invariant, then showing that this property is a valid loop invariant will be equivalent to proving correctness of the algorithm. To show that the loop invariant is correct we must show:

Initialization: the property is true before the first iteration of the loop.

Maintenance: if the property is true before a given iteration, it is also true before the next iteration.

Termination: The loop terminates and gives a useful property that helps us to verify that the algorithm is correct.

For example, insertion sort is correct if the final array is sorted. Consider the following loop invariant:

At the start of each iteration of the for loop, the subarray $A[1:i-1]$ consists of the elements originally in $A[1:i-1]$, but in correct sorted order.

To prove correctness of insertion sort, we prove the loop invariant always holds by induction:

Initialization: Before the first loop iteration¹, $i = 2$ so the left subarray is just $A[1:i-1] = A[1]$, which clearly contains the original element and is sorted.

Maintenance: Formally, we should check the while loop with a loop invariant. Informally, each iteration moves the values by one until $A[i]$ is placed in the correct position. This implies that $A[1:i]$ is still sorted and contains the same elements. Thus, *incrementing* i preserves this.

Termination: To conclude our proof, we want to take the condition that terminates the loop and sub it into the original formulation of the loop invariant above. This should give us a statement that proves the correctness of the algorithm. The loop stops where $i > n$, or equivalently, $i = n + 1$. Subbing this in gives the following statement: *The subarray $A[1:n]$ consists of elements originally in $A[1:n]$ but in sorted order.* This is exactly our goal (that is, array A has been sorted) and so this proves correctness by induction.

Linear search solves the following *searching problem*:

Input: A sequence of n numbers $\langle a_1, a_2, \dots, a_n \rangle$ stored in array $A[1:n]$ and a value x .

Output: An index i such that x equals $A[i]$ or the special value NIL if x does not appear in A .

Linear search scans the array from beginning to end, looking for x :

```

LINEAR-SEARCH(A, n, x)
  for i in 1 to n
    if A[i] = x
      return i
  return NIL

```

Prove its correctness with the following loop invariant: *At the start of each iteration, the subarray $A[1:i-1]$ contains only elements not equal to x .* This is valid as the loop terminates when returning either i or NIL.

2 Analysis of Algorithms

Analysing an algorithm has come to mean predicting the resources that it requires, most commonly time (and memory). In the *random-access machine* (RAM) model of computation,

¹In a for loop, this is between the counter initialization `int i = 2` and the condition `i ≤ n` assuming a C for loop: `for (int i = 2; i ≤ n; i++)`.

each instruction or data access takes a constant amount of time, and no operations occur concurrently. Note that we can't assume that every instruction takes constant time as this would lead to an unrealistic and unhelpful model. For example, an instruction that sorts in a single instruction, would be assumed to take constant time in the RAM model, which is unrealistic. Thus, the only computations that we consider to take constant time are those that can occur in real computers, especially arithmetic (standard operations, mod, floor, ceiling), data management (load, store, copy) and control (conditional & unconditional branch, subroutine call, return).

The notion of input size depends on the problem being studied. For example, in sorting, the number of items in the input is the clear choice; when multiplying two integers, the total number of bits needed to represent the input in binary may be used; for graph algorithms, both the number of vertices and edges are useful input sizes. The running time of an algorithm depends on a particular input, and is the number of instructions and data accesses executed. Assume each line of pseudocode takes constant time. Note that, while calling a subroutine takes a constant amount of time, actually executing this subroutine may (and likely will) take more.

We shall analyse insertion sort to demonstrate this analysis process. See Figure 2 for the pseudocode of the algorithm with the cost of each step. Denote the number of times the while loop executes the **for** $i = 2$ to n with t_i . Thus, we must sum over the i to get the number of times these lines will execute. Note that the **for** and **while** loop execute one extra time than their bodies to account for the terminating check (that is, they execute $(n - 1) + 1$ times). The running time is the cost multiplied by the number of times this step occurs.

INSERTION-SORT(A, n)	<i>cost</i>	<i>times</i>
1 for $i = 2$ to n	c_1	n
2 $key = A[i]$	c_2	$n - 1$
3 // Insert $A[i]$ into the sorted subarray $A[1 : i - 1]$.	0	$n - 1$
4 $j = i - 1$	c_4	$n - 1$
5 while $j > 0$ and $A[j] > key$	c_5	$\sum_{i=2}^n t_i$
6 $A[j + 1] = A[j]$	c_6	$\sum_{i=2}^n (t_i - 1)$
7 $j = j - 1$	c_7	$\sum_{i=2}^n (t_i - 1)$
8 $A[j + 1] = key$	c_8	$n - 1$

Figure 2: Pseudocode of insertion sort with the cost of each step.

Thus the running time for insertion sort is $T(n) = c_1 + c_2(n - 1) + c_4(n - 1) + c_5 \sum_{i=2}^n t_i + c_6 \sum_{i=2}^n (t_i - 1) + c_7 \sum_{i=2}^n (t_i - 1) + c_8(n - 1)$. Note that, while defining t_i in this way simplifies the function, we could have defined it in a different way if that were more useful.

In the best case scenario, the array was already sorted, so each **while** loop is only executed once. That is, $t_i = 1$ (for all i) so $T(n) = (c_1 + c_2 + c_4 + c_5 + c_8)n - (c_2 + c_4 + c_5 + c_8) = an + b$, for some constants a and b . This is a *linear* function of n .

In the worst case scenario, the array is in reverse sorted order initially. In this case, $t_i = i$ as $A[j] > key$ is true every time. We can prove the following relations by induction.

$$\sum_{i=2}^n i = \left(\sum_{i=1}^n i \right) - 1 = \frac{n(n+1)}{2} - 1$$

$$\sum_{i=2}^n (i - 1) = \left(\sum_{i=1}^{n-1} i \right) = \frac{n(n-1)}{2}$$

If we write out $T(n)$ again, we get a function of the form $T(n) = an^2 + bn + c$, where a , b and c depend on the constants c_k . This is a *quadratic* function of n .

Generally, we concentrate on the worst-case running time as this gives us an upper bound for the running time of our algorithm. Average-case running time can also be useful (for example, in the above case $t_i = \frac{i}{2}$), although it's often the same as the worst-case. We are interested in rate of growth (or order of growth) of the running time, that is, only the leading term. We write the order of growth with theta notation. For example, insertion sort has $\Theta(n^2)$.

It is not common to consider the best-case running time, although there are reasons to do so. For example, the best case of insertion sort is linear time (which is the fastest conceivable comparison sort time as every element needs to be "looked at" at least once). We can use this initially to check if the array is already sorted, before applying some other sorting algorithm if it isn't. This is no more computationally expensive than just running the second algorithm immediately, which will have a running time that is greater than linear for the worst case. Thus, we can get the benefit of quickly sorting the special case for practically no cost.

2.1 Selection Sort

We now introduce another sorting algorithm: *selection sort*. This maintains the following loop invariant: *at the start of each outer for loop, the subarray $A[1:i - 1]$ consists of the $i - 1$ smallest elements in the array $A[1:n]$, and they are in sorted order*. It starts with a left (empty) subarray and the right subarray (which is just the array). It simply searches the right subarray for the lowest element and exchanges it with the first element in the right subarray. It then iterates over this, considering the left subarray to contain the former first element of the right array, thus shrinking the right array. This is made more explicit in the pseudocode below:

```

SELECTION-SORT(A, n)
  for i = 1 to n - 1
    smallest = i
    for j = i + 1 to n
      if A[j] < A[smallest]
        smallest = j
    exchange A[i] with A[smallest]

```

This has average-case and worst-case running time $\Theta(n^2)$.

2.2 Bubble Sort

We will briefly mention another simple but inefficient (that is, $\Theta(n^2)$) sorting algorithm.

```

BUBBLE-SORT(A, n)
  for i = 1 to n - 1
    for j = n downto i + 1
      if A[j] < A[j - 1]
        exchange A[j] with A[j - 1]

```

3 Merge Sort

There are many algorithm design techniques. Insertion sort used the incremental method. We will now consider the *Divide-and-conquer method*.

Many useful algorithms are recursive in structure: to solve a given problem, they recurse (call themselves) one or more times to handle closely related subproblems. These algorithms typically follow the divide-and-conquer method: if the problem is small enough (as in the

base case), you solve directly without recursing. Otherwise, you perform the following steps: **divide** the problem into smaller instances of the same problem; **conquer** the subproblems by solving recursively; **combine** the subproblem solutions to form a solution to the original problem.

The merge sort algorithm follows divide-and-conquer closely. **Divide** the subarray to be sorted, $A[p:r]$, into two adjacent subarrays, each of half the size (approximately). Do this by computing the midpoint q of $A[p:r]$ (taking the average of p and r) and thus dividing A into subarrays $A[p:q]$ and $A[q + 1:r]$. **Conquer** by sorting each of the two subarrays recursively using merge sort. **Combine** by *merging* the sorted subarrays back into a single sorted subarray, producing the sorted answer. The base case is when the subarray is just one element (that is, when $p = r$).

```

MERGE(A, p, q, r)
 $n_L = q - p + 1$     // Length of first subarray
 $n_R = r - q$         // Length of second subarray
let  $L[0:n_L - 1]$  and  $R[0:n_R - 1]$  be new arrays

for  $i = 0$  to  $n_L - 1$ 
     $L[i] = A[p + i]$ 
for  $j = 0$  to  $n_R - 1$ 
     $R[j] = A[q + 1 + j]$ 

 $i = 0$ 
 $j = 0$ 
 $k = p$ 

while  $i < n_L$  and  $j < n_R$ 
    if  $L[i] \leq R[j]$ 
         $A[k] = L[i]$ 
         $i = i + 1$ 
    else
         $A[k] = R[j]$ 
         $j = j + 1$ 
     $k = k + 1$ 

// Copy remainder of unmerged array across
while  $i < n_L$ 
     $A[k] = L[i]$ 
     $i = i + 1$ 
     $k = k + 1$ 
while  $j < n_R$ 
     $A[k] = R[j]$ 
     $j = j + 1$ 
     $k = k + 1$ 

```

It may be easier to visualise this as two sorted piles of cards face up (that is so only 2 cards are visible at a time) with an empty stack face down. At each iteration, take the visible card with the lowest value and place it face down of the initially empty pile. Iterate until one of the face up piles is empty. At this point, flip the remaining pile and add it to the end of the facedown stack. Note that, as both piles are sorted, the best case takes $n / 2$ steps as it is necessary to individually copy every card in at least one of the piles. The worst case is $n - 1$ (that is, going through every card). Thus merging takes $\Theta(n)$ time. To see this, observe

that the variable initializations take constant time, the for loops take $\Theta(n_L + n_R) = \Theta(n)$ and the outcome of all the while loops is to copy every element in the array across once. Thus, the running time is $\Theta(n)$.

```

MERGE-SORT(A, p, r)
    if p >= r          // zero or one elements
        return
    q = ⌊(p + r) / 2⌋
    MERGE-SORT(A, p, q)
    MERGE-SORT(A, q + 1, r)
    MERGE(A, p, q, r)

```

Note that $\lceil x \rceil$ denotes the least integer that is greater than or equal to x , and $\lfloor x \rfloor$ denotes the greatest integer that is less than or equal to x .

As an aside, *coarsening* the leaves of the recursion and using insertion sort for subproblems of a specific (and sufficiently small size) can sometimes produce a faster running time. This splits the array into $\frac{n}{k}$ subarrays of size k (instead of splitting the array into n subarrays of size 1 as in traditional merge sort) and sorts each subarray with insertion sort, then merging them as normal. This modified algorithm has running time $\Theta(nk + n \lg(n/k))$.

3.1 Analysis of Divide-and-Conquer Algorithms

When an algorithm contains a recursive call, you can often describe its running time by a *recurrence equation*, which describes the overall running time on a problem of size n , in terms of the running time of the same algorithm with smaller inputs.

Generally, let the worst case running time of a divide-and-conquer algorithm be given by

$$T(n) = \begin{cases} \Theta(n) & \text{if } n < n_0 \\ D(n) + aT(\frac{n}{b}) + C(n) & \text{otherwise} \end{cases}$$

where $D(n)$ is the cost to divide and $C(n)$ is the cost to combine, and the problem is divided into a subproblems of size $\frac{n}{b}$. For example, merge sort divides into two subproblems of size $\lfloor \frac{n}{2} \rfloor$ and $\lceil \frac{n}{2} \rceil$ respectively.

For merge sort, the divide step computes the middle of the subarray, which is constant time. The conquer step recursively solves two subproblems of size $\frac{n}{2}$ so contributes a $2T(\frac{n}{2})$ term. Combine is executed by merge, which we showed takes linear time. Thus $T(n) = 2T(\frac{n}{2}) + \Theta(n)$, which we can also write as

$$T(n) = \begin{cases} c_1 & \text{if } n = 1 \\ 2T(\frac{n}{2}) + c_2n & \text{if } n > 1 \end{cases}$$

It turns out that $T(n) = \Theta(n \lg n)$, which we shall prove later. This can be intuitively seen from a recursion tree (see Figure 3).

3.2 Binary Search

Another example of a divide-and-conquer algorithm is *binary search*, which is a searching algorithm. It specifically searches for some element x in a sorted array A . It does this by finding the midpoint of the boundaries of the array being searched and compares the value in this position to x . If they are equal, it returns the midpoint (specifically, the index). If $x > A[mid]$, it performs the same action on the right (higher) half of the array. Otherwise, it performs the same action on the left (lower) half of the array. This can be done iteratively

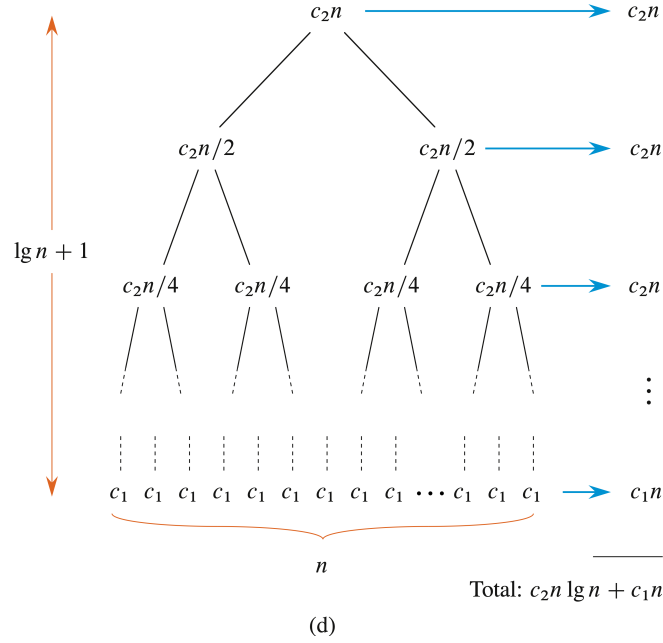



Figure 3: Recursion tree providing intuition for the total cost, by progressively expanding $T(n)$ until it reaches the base case layer (the leaves).

(by using a while loop that checks if the boundaries have crossed each other and updating the boundaries after each iteration) or recursively (by calling binary search on progressively smaller subarrays). After each iteration, the number of elements to be checked halves, so binary search has a recurrence $T(n) = T(\frac{n}{2}) + \Theta(1)$. Thus, the worst-case running time is $\Theta(\lg n)$, as this solves the recurrence.

```

RECURSIVE-BINARY-SEARCH(A, x, low, high)
  if low > high
    return NIL
  mid = ⌊(low + high) / 2⌋
  if x == A[mid]
    return mid
  else if x > A[mid]
    return RECURSIVE-BINARY-SEARCH(A, x, mid + 1, high)
  else
    return RECURSIVE-BINARY-SEARCH(A, x, low, mid - 1)

```

4 Complexity and O-notation

When we look at input sizes sufficiently large such that the order of growth is the only thing that impacts the running time, we are studying the *asymptotic* efficiency of algorithms. That is, how does the running time of an algorithm depend on the size of the input n in the limit $n \rightarrow \infty$.

4.1 Asymptotic Notation

O-notation: upper bound on the asymptotic behaviour of a function. That is, the function grows no faster than a certain rate (based on the highest order term).

Ω -notation: lower bound on the asymptotic behaviour of a function. That is, the function grows at least as fast as a certain rate (based on the highest order term).

Θ -notation: tight bound on the asymptotic behaviour of a function. That is, the function grows precisely at a certain rate (based on the highest order term).

Note that saying the worst case running time is $\Omega(n^2)$ is not the same as saying that all inputs take at least cn^2 steps, but rather that there exists at least one input, of size n , that takes at least cn^2 for every input size n above a certain value.

For example, insertion sort has $O(n^2)$ as the outer **for** loop undergoes $n - 1$ iterations and the worst case for the inner **while** loop is for it to iterate $i - 1$ times (with $i = n$ as it is the worst case). Thus, the running time is characterised by $(n - 1)(n - 1) = O(n^2)$. It also has $\Omega(n^2)$ by the following argument: assume that the input array has size n and that the first αn values are the largest, where $0 < \alpha \leq \frac{1}{2}$. To sort by insertion sort, these αn values need to be moved through the middle $(1 - 2\alpha)n$ positions, one position at a time, to end up in the last αn positions. Thus, there are at least $\alpha(1 - 2\alpha)n^2 = \Omega(n^2)$ steps². Given that the lower and upper bounds both constrain insertion sort to quadratic time, we can also say that it has a running time of $\Theta(n^2)$.

²As an aside, this value is maximised when $\alpha = \frac{1}{4}$